

SATURN AI

Private Personal AI Assistant

Research Document & Implementation Blueprint

Prepared for:

Chris — New Era Designs, LLC

March 2026

This document is a complete, practical blueprint for building Saturn AI — a private, local-first AI assistant designed for a solo design business owner. It covers architecture, hardware, software stack, automation workflows, phased implementation, and a first-week setup checklist. No cloud dependency required. No hype. Just a buildable system.

CONFIDENTIAL — FOR PERSONAL USE ONLY

01

Executive Summary

What Is Saturn AI?

Saturn AI is your private, always-available AI assistant built specifically for New Era Designs, LLC. It runs on your existing Windows workstation and knows your business, files, clients, workflow, and voice. It replaces daily friction: hunting for files, drafting proposals from scratch, summarizing project notes, tracking deadlines, and transcribing voice ideas into structured tasks.

The Three Versions of Saturn AI

Version	Name	What It Does	Timeline	Complexity
V1	Saturn Core	Local chat + document Q&A; Ask questions, draft proposals, search knowledge base. Text-only via browser interface.	Month 1–2	Low–Medium
V2	Saturn Voice	V1 plus voice input (speak to Saturn), voice output (Saturn replies aloud), and automations via n8n.	Month 3–5	Medium
V3	Saturn Ops	Full business ops: CRM, invoice reminders, proposal tracking, intake automation, multi-device access.	Month 6–12	Medium–High

What Is Realistic in 2026 for One Person?

Version 1 is buildable in a weekend. Versions 2 and 3 require a few evenings each over several weeks. You do not need to write code. You need patience, a willingness to configure software, and the ability to follow structured instructions.

Realistic expectation: Saturn AI will save you 1–3 hours per day in admin, drafting, and file hunting tasks once fully set up. The cleaner your file system, the smarter Saturn becomes. Your RTX 4060 workstation already exceeds minimum requirements for Version 1 and 2.

02 Primary Use Case Definition

Must-Have First Features (Phase 1)

These use cases have the highest ROI for a solo designer and require the least setup.

Use Case	What It Solves	ROI	Ease	Phase
Document Q&A;	Ask Saturn: "What did I charge for the Rubicon project?" and get an answer from your own files.	Very High	Easy	1
Proposal Drafting	Saturn pulls your templates and past proposals to help you write new ones in minutes.	Very High	Easy	1
File Search	Natural language retrieval: "Find the Rubicon sign schedule" instead of folder-diving.	High	Easy	1
Naming Convention Enforcer	Ask Saturn the correct file naming format. Paste a file list, get corrected names back.	High	Easy	1
Meeting Notes to Tasks	Paste raw notes, get a clean bullet-point task list with owners and deadlines.	High	Medium	1–2
Pricing & Scope Templates	"Build a scope estimate for a 3-sign wayfinding system for a gym." Saturn uses your pricing logic.	High	Easy	1
Creative Brief Summarizer	Paste a client intake form. Saturn extracts deliverables, deadlines, and scope.	High	Easy	1

Cool Later Features (Phase 2–3)

Use Case	What It Adds	Phase
Voice Input / Output	Speak to Saturn hands-free while working in SketchUp or Illustrator. Saturn reads responses aloud.	2
Voice-to-Task Capture	Hit a hotkey, speak a task, Saturn transcribes and adds it to your task list in Notion.	2
CRM Follow-Up Reminders	Saturn scans a contact log and reminds you when a client hasn't been contacted in 14+ days.	2
Invoice Follow-Up Automation	Trigger an automated email draft when an invoice passes a set number of days unpaid.	2–3
Weekly Business Review	Every Sunday, Saturn auto-generates a summary: completed tasks, open projects, upcoming deadlines.	2
Project Folder Auto-Creator	Type "New project: Pulse Gym Phase 1" and Saturn creates your full folder tree automatically.	2
Email Drafting Assistant	Paste an email; Saturn drafts a professional response in your voice and tone.	3
Proposal Version Tracker	Saturn logs proposal revisions and reminds you when a proposal goes stale.	3

03

System Architecture Options

The Three Architecture Models

Architecture	Privacy	Cost	Performance	Verdict
Local-Only	Maximum — nothing leaves PC	\$0/mo after hardware	Good — 7B–13B models on RTX 4060	Private but limited writing quality
Hybrid Local + Cloud API (Recommended)	High — files stay local, only typed prompts leave	\$5–20/mo API	Excellent — best of both worlds	Best balance for Saturn V1 and V2
Cloud-Only (e.g., ChatGPT API)	Low — files and prompts go to third party	\$20–60/mo	Excellent	Do NOT use for client files

Recommendation: Hybrid Local + Cloud API. Run Ollama locally for private document Q&A.; Use OpenAI GPT-4.1 mini for complex writing. Your files NEVER leave your machine — only your typed prompt does.

Architecture Data Flow

- [YOU] speak or type [AnythingLLM interface in Chrome]
- ↓ connects to [Ollama: Mistral 7B running on your RTX 4060 GPU]
- ↓ for complex writing [OpenAI GPT-4.1 mini API — prompt only, no files]
- ↓ document search [ChromaDB: indexes your _SATURN_KB/ folder locally]
- ↓ automations [n8n: project folders, reminders, weekly review]

Tradeoff Comparison

Concern	Local-Only	Hybrid (Recommended)	Cloud-Only
Client files leave PC?	Never	No — only typed prompt	Yes — risky
Internet required?	No	Only for complex drafting	Always
Monthly cost?	\$0	\$5–20/mo	\$20–60/mo
Response quality?	Good	Excellent	Excellent
Setup difficulty?	Medium	Medium	Easy but least private
Works offline?	Yes	Partially	No

04

Hardware Requirements

Your current RTX 4060 workstation already exceeds minimum requirements for Saturn V1 and V2. Start building today — zero hardware purchases needed for Phase 1.

Hardware Requirements Table

Category	Entry-Level	Mid-Tier (Your Level)	Advanced Option	Price Range	Why It Matters
GPU	RTX 3060 12 GB	RTX 4060 8 GB	RTX 4090 24 GB	\$200–1,600	VRAM determines model size. More VRAM = smarter local responses.
CPU	i5-12400 / R5 5600	i7-14700 / R7 5700X	R9 7950X / i9-14900K	\$150–450	CPU handles embedding, n8n automation, and multitasking.
System RAM	16 GB DDR4	32 GB DDR4/DDR5	64 GB DDR5	\$40–200	Running Windows + design apps + Ollama + n8n needs 32 GB minimum.
Primary SSD	500 GB NVMe	1 TB NVMe SSD	2 TB NVMe SSD	\$60–150	Models are 4–8 GB each. Fast SSD required for model loading speed.
Secondary Storage	2 TB external HDD	4 TB external HDD	8 TB NAS Synology DS223	\$60–400	Knowledge base files, backups, and project archives.
Microphone	Built-in (poor)	HyperX SoloCast \$60	Shure MV7+ \$200	\$0–250	Mic quality directly affects Whisper transcription accuracy.
UPS Power	APC BE600M1 \$70	APC BR1500MS2 \$200	CyberPower CP1500 \$250	\$70–250	Prevents data corruption if power cuts during model runs.
Network/Router	Existing ISP router	Wi-Fi 6 TP-Link AX3000	Wi-Fi 6E + LAN switch	\$0–200	Only relevant if adding a separate server in Version 3.

GPU VRAM Guide for Local Models

GPU VRAM	Best Model to Run	Quality	Use Case for Saturn
8 GB (Your RTX 4060)	Mistral 7B Q4, Llama 3.1 8B Q4	Good	Solid for doc Q&A, proposal drafting, summaries. Ideal for Saturn V1.

GPU VRAM	Best Model to Run	Quality	Use Case for Saturn
12 GB (RTX 3060)	Mistral 7B Q8, Llama 3.1 8B, Gemma 2 9B	Very Good	Higher quality responses. Recommended upgrade path.
16–24 GB (RTX 3090/4090)	Mistral 22B, Llama 3.1 70B Q4, Qwen2.5 32B	Excellent	Near-GPT-4 quality locally. Best long-term target.

05

Software Stack

Complete Saturn AI Software Stack

Tool	Purpose	OS	Cost	Difficulty	Essential?
Ollama	Runs local AI models (Mistral, Llama, Gemma) on your GPU. Command-line model manager.	Win/Mac/Linux	Free	Easy	YES — Core
LM Studio	GUI-based alternative to Ollama. Download and run models visually. Great for beginners.	Win/Mac/Linux	Free	Very Easy	Optional
AnythingLLM	All-in-one chat + RAG interface. One-click Windows installer. Best starting point.	Win/Mac/Linux	Free	Easy	YES — Start Here
Open WebUI	ChatGPT-like browser UI for Ollama. More features but requires Docker.	Windows (Docker)	Free	Medium	Optional (Phase 2)
n8n	Visual workflow automation (like Zapier, self-hosted). Connects Saturn to Notion, Gmail, folders.	Win/Mac/Linux	Free self-host / \$24 cloud	Medium	YES — Phase 2
ChromaDB	Local vector database. Saturn searches your documents through this.	Windows (Docker/Python)	Free	Medium	YES — Built into AnythingLLM
Qdrant	More advanced local vector DB. Better for large document collections (1,000+ files).	Windows (Docker)	Free	Medium-High	Optional (scale later)
Whisper (faster-whisper)	Local speech-to-text. Converts your spoken words to text for Saturn.	Win/Linux	Free	Medium	YES — Phase 2
Kokoro TTS	Local text-to-speech. Saturn speaks responses in a natural voice. 54 voices.	Windows (Docker)	Free	Medium	YES — Phase 2
Piper TTS	Lighter-weight local TTS. Simpler than Kokoro, fewer voices.	Win/Linux	Free	Easy	Optional alt to Kokoro
Docker Desktop	Runs apps in isolated containers. Required for Open WebUI, n8n server, Qdrant.	Windows (WSL2)	Free personal	Medium	YES — Phase 2
Everything (voidtools)	Instant Windows file name search across all drives.	Windows only	Free	Very Easy	YES — Install today
DocFetcher	Full-text search inside PDFs and DOCX files on your local drives.	Win/Mac/Linux	Free	Easy	YES — Phase 1
Obsidian	Local-first markdown notes. Knowledge base for SOPs, naming rules, templates.	Win/Mac/Linux	Free / \$25 sync	Easy-Medium	Recommended — Phase 1
OpenAI API	Cloud AI fallback. GPT-4.1 mini for complex writing. ~\$0.15/1M input tokens.	Cloud API (prompt only)	\$5–20/mo	Easy	YES — Hybrid setup
Anthropic (Claude)	Alternative cloud writing fallback. Claude Haiku is excellent and cheap.	Cloud API (prompt only)	\$5–15/mo	Easy	Optional alt
Mistral 7B (Ollama)	Best all-around 7B local model for Q&A, writing, reasoning on 8 GB VRAM.	Via Ollama	Free	1 command	YES — V1 model

Tool	Purpose	OS	Cost	Difficulty	Essential?
nomic-embed-text	Embedding model. Converts documents to searchable vectors. Required for RAG.	Via Ollama	Free	1 command	YES — Required

Minimum V1 Stack — Install These 5 Things

1. Ollama (local model runner) — ollama.com
 2. AnythingLLM (chat interface + RAG) — anythingllm.com
 3. Mistral 7B via Ollama: `run ollama pull mistral`
 4. nomic-embed-text via Ollama: `run ollama pull nomic-embed-text`
 5. Everything by voidtools (instant file name search) — voidtools.com
- Optional but recommended immediately: DocFetcher (full-text file search) and Obsidian (knowledge base notes).

06

OS & Platform Guidance

Should Saturn Run on Windows, Linux, Docker, or a Server?

Approach	What It Means	Verdict for You
Native Windows	Install all tools directly on Windows 11 Pro. Ollama, AnythingLLM, n8n desktop — all install like normal apps.	Start here for V1. Best for non-programmers. No barriers.
WSL2 (Windows Subsystem for Linux)	Run Linux inside Windows. Most AI tools work natively. Better Docker compatibility.	Enable WSL2 for Phase 2 when you add Docker tools.
Docker Desktop on Windows + WSL2	Run apps in isolated containers. Open WebUI, n8n server, Qdrant, Kokoro TTS all run as containers.	Use for Phase 2 components. AnythingLLM and Ollama stay native.
Dedicated Mini Server	A separate PC runs Saturn 24/7 independently from your workstation.	Consider for Version 3 only. Not needed until then.
Full Linux Migration	Move entirely to Ubuntu/Fedora.	Do NOT do this. You would lose Adobe CC, ECDesign, and SketchUp.

For the first 2–3 months, everything lives on your main Windows PC with no Docker required. Start native, add Docker only when you need Open WebUI or Kokoro TTS in Phase 2.

What Should Live Where?

Component	Where It Lives	Why
Ollama (model runner)	Main Windows PC — native install	Uses GPU directly. Native Windows version works well.
AnythingLLM	Main Windows PC — native install	One-click Windows installer. No Docker needed.
AI Models (Mistral, Llama)	Main PC — SSD drive	Models need fast SSD. Keep on same machine as Ollama.
ChromaDB (vector database)	Inside AnythingLLM — automatic	AnythingLLM manages this internally. Zero setup.
n8n automation	Main PC via Docker or n8n desktop app	Desktop app simpler for Phase 1. Docker for Phase 2.
Kokoro TTS / Whisper	Main PC via Docker (Phase 2)	GPU-accelerated. Runs as Docker container.
Knowledge Base files	External drive or secondary SSD	Keep separate from OS drive. Easier to back up.
Notion / Google Workspace	Cloud — unchanged	These stay cloud. n8n connects to them via API.
Backups	External HDD + Google Drive	Local weekly backup + cloud for catastrophic loss.

07

Knowledge Base & Memory System

File Type Compatibility

File Type	Can Saturn Read It?	How to Feed It	Notes
PDF	Yes — excellent	Drag into AnythingLLM workspace	Proposals, contracts, spec sheets.
Word (.docx)	Yes	Drag into AnythingLLM workspace	SOPs, meeting notes, pricing templates.
Text (.txt, .md)	Yes — best format	Drag or sync folder	Plain text notes, naming rules, process docs.
Excel/Sheets (.xlsx, .csv)	Partial — CSVs work best	Export to CSV first	Pricing sheets, client lists, sign schedules.
Email exports (.eml)	Yes if exported as .txt	Export from Outlook as .txt	Client correspondence, project threads.
Images (.jpg, .png)	No — V1 limitation	Not applicable	Saturn cannot read text inside images.
SketchUp (.skp)	No	Export notes as PDF or TXT separately	Write a companion text file per project.
Notion pages	Yes (export to PDF/Markdown)	Export from Notion, import to Saturn	Or use n8n to auto-sync Notion to text files.

Recommended Folder Structure — _SATURN_KB/

Saturn's intelligence comes entirely from what you feed it. Only clean, current, well-named files go into this folder.

Folder	Contents	Naming Convention
_SATURN_KB/	Root folder for all AI-indexed files only	Do not mix with active project files
_SATURN_KB/00_Business/	LLC docs, pricing philosophy, rate cards, SOPs	NERA_SOP_[Topic]_v[N].pdf
_SATURN_KB/01_Proposals/	All past proposals as PDFs	NERA_PROP_[ClientCode]_[Phase]_v[N].pdf
_SATURN_KB/02_Pricing/	Pricing sheets, scope templates, unit cost CSVs	NERA_PRICE_[Category]_[Year].csv
_SATURN_KB/03_Projects/	One subfolder per project + a project brief TXT	NERA_PROJ_[ClientCode]_BRIEF.txt
_SATURN_KB/04_Clients/	Client profiles, contact logs, intake forms	NERA_CRM_[ClientCode]_PROFILE.txt
_SATURN_KB/05_Templates/	Proposal, email, brief templates	NERA_TMPL_[Type]_[Name].docx
_SATURN_KB/06_Naming_Rules/	Your file naming conventions in plain text	NERA_NAMING_RULES.txt
_SATURN_KB/07_Meetings/	Meeting notes as TXT or PDF	NERA_MTG_[ClientCode]_[YYYYMMDD].txt

Key Risks of Messy File Systems

- Duplicate files cause Saturn to return conflicting information. One source of truth per document.
- Old proposal versions quoted as current pricing. Archive outdated files before indexing.
- Scanned PDFs (image-only) cannot be read. Run OCR first via Adobe Acrobat or free PDF24.
- Files over 50 MB may cause AnythingLLM to time out. Split large documents into sections.
- Never feed Saturn your entire Downloads folder or desktop — only the _SATURN_KB/ folder.

08

Automation Tasks Saturn Could Handle

Task	Example Workflow	Tools	Impact	Phase
New project folder creation	Type "New project: Pulse Gym Phase 1" 📄 n8n creates full folder tree using your naming convention.	n8n + Python script	High	2
Naming convention enforcement	Paste a file list 📄 Saturn returns corrected names based on NERA naming rules in knowledge base.	AnythingLLM only	High	1
Project summary generation	"Summarize the Rubicon project" 📄 Saturn reads project brief and proposal, returns a 1-paragraph summary.	AnythingLLM	High	1
Voice note to task list	Press hotkey, speak: "Call Marcus about sign specs by Friday" 📄 Saturn transcribes into structured task.	Whisper + AnythingLLM + n8n	High	2
Proposal template pull	"Start a proposal for a 5-sign retail system at \$85/hr" 📄 Saturn fills template with standard language.	AnythingLLM + template DOCX	Very High	1–2
CRM follow-up reminders	n8n checks client CSV daily. If a contact hasn't been emailed in 14+ days, sends a Notion reminder.	n8n + Google Sheets/CSV	High	2
Weekly business review	Every Sunday at 6pm, n8n pulls open projects and tasks. Saturn formats into a weekly brief.	n8n + Notion API + AnythingLLM	High	2–3
Intake form data extraction	Paste a client intake response 📄 Saturn extracts: client name, project type, budget, timeline, deliverables.	AnythingLLM only	Medium	1
Downloaded file organizer	n8n watches Downloads folder. New PDFs/images move to correct project folder by filename pattern.	n8n + file watcher	Medium	2
Email drafting support	Paste a client email 📄 "Draft a professional response declining the timeline." Saturn drafts in your voice.	AnythingLLM + email notes	High	1
Invoice follow-up	n8n checks invoice tracker CSV. If invoice is 7+ days past due, drafts polite follow-up to your Drafts.	n8n + Sheets + Gmail API	Very High	2–3
Proposal version tracking	Every time you save a proposal PDF, n8n logs filename, date, client, and version to a Google Sheet.	n8n + Google Sheets	Medium	2
Creative brief summarization	Paste a full brief 📄 Saturn extracts: objective, audience, deliverables, brand constraints, deadline.	AnythingLLM only	High	1
Meeting notes formatting	Paste raw notes 📄 Saturn produces: decisions made, action items with owners and due dates, open questions.	AnythingLLM only	Very High	1

09

Pricing & Budgeting

Tier 1 — Bare Minimum Starter (Recommended to Start)

Best for: Testing Saturn immediately at zero cost using your existing hardware.

Item	Cost	Notes
All software (Ollama, AnythingLLM, Everything, DocFetcher)	\$0	100% free and open source.
AI models (Mistral 7B, nomic-embed-text)	\$0	Downloaded free via Ollama.
OpenAI API (optional cloud writing fallback)	\$5–15/mo	Pay-as-you-go. Under \$15/mo for solo use.
New hardware	\$0	Your RTX 4060 workstation handles everything.
TOTAL One-Time Cost	\$0	You already have what you need.
TOTAL Monthly Cost	\$0–15/mo	Optional cloud API only.

Tier 2 — Practical Small Business Setup

Best for: A fully functional Saturn V1 + V2 with voice and automation.

Item	Type	Cost	Notes
HyperX SoloCast or Blue Yeti Nano microphone	One-time	\$50–70	Essential for Phase 2 voice input quality.
APC BE600M1 UPS power protection	One-time	\$70	Protects workstation during model runs.
External 4 TB backup drive	One-time	\$80–100	Knowledge base backups and project archive.
OpenAI API — GPT-4.1 mini	Monthly	\$10–20/mo	Excellent writing quality at very low cost.
n8n self-hosted via Docker	One-time setup	\$0	Free to self-host. One evening to configure.
Anthropic Claude Haiku (optional)	Monthly	\$5–10/mo	Secondary AI for drafting. Very affordable.
TOTAL One-Time Hardware		\$200–240	Mic + UPS + drive.
TOTAL Monthly Software		\$15–30/mo	OpenAI API + optional Claude.

Tier 3 — Premium Prosumer Saturn

Best for: Long-term serious use, larger local models, always-on server, NAS storage.

Item	Type	Cost	Notes
GPU Upgrade: RTX 3090 24 GB (used) or RTX 4080 16 GB	One-time	\$500–900	Enables 22B–70B models. Near GPT-4 quality locally.
Dedicated mini server: Beelink SER7 or Intel NUC 13 Pro	One-time	\$350–600	Runs Saturn 24/7 off your main workstation.
Synology DS223 NAS + 2x 4 TB drives	One-time	\$400–500	Network storage with automatic backups.
Shure MV7+ or Elgato Wave:3 microphone	One-time	\$150–200	Broadcast quality for best Whisper accuracy.
n8n Cloud plan	Monthly	\$24/mo	Optional — managed and accessible from anywhere.
OpenAI API GPT-5 Mini for premium tasks	Monthly	\$15–40/mo	Access to frontier models for complex proposals.
TOTAL One-Time		\$1,650–2,450	GPU + server + NAS + mic + UPS.
TOTAL Monthly		\$40–65/mo	n8n Cloud + API usage.

Hidden Costs to Know About

Hidden Cost	Estimated	How to Minimize
API overuse from pasting large documents to cloud	\$5–30 spike	Never send full files to cloud API. Use local model for doc Q&A.;
Time for file cleanup before starting	5–15 hours	Budget a weekend. One-time cost that pays off immediately.
Electricity from GPU when running models	\$5–15/mo	Configure Windows to idle GPU when Ollama is not active.
SSD space for multiple AI models	\$0 (space)	Each model = 4–8 GB. Budget ~50 GB for 5–6 models.
Learning time across all phases	20–40 hours total	Follow the phased plan in Section 11. Do not rush.

10

Security & Privacy

Core Privacy Principles

Principle	How to Apply It	Risk if Ignored
Client files never leave your PC	All document Q&A; uses local Ollama + local ChromaDB. Never paste client contracts into cloud APIs.	Client data could be logged by OpenAI/Anthropic and used for training.
Use cloud API for writing only, not reading	Use OpenAI/Claude only for text generation from your typed prompts — not for analyzing uploaded files.	Uploading a client PDF to OpenAI exposes it to third-party servers.
Opt out of API training data	OpenAI API Settings  Privacy  disable "Improve our models for everyone." Anthropic doesn't train on API data by default.	Your prompts could be used to train models without your consent.
Encrypt external backup drive	Use Windows BitLocker (free, built-in) on your external knowledge base drive.	A stolen drive exposes all client data, proposals, and pricing.
2FA on all API accounts	Enable two-factor authentication on OpenAI, Anthropic, n8n Cloud, and Notion.	Compromised API account = unauthorized charges and data exposure.
Only index _SATURN_KB/ folder	Never give AnythingLLM access to your entire C: drive or Downloads folder.	Saturn surfaces irrelevant or personal data in responses.
Keep tools updated monthly	Update Ollama, AnythingLLM, and Docker once a month.	Unpatched tools can have exploitable vulnerabilities on your network.

Backup Strategy (4-Layer)

Layer	What to Back Up	Frequency	Tool
External Drive	_SATURN_KB/ + all project files	Weekly (automated)	FreeFileSync (free)
Cloud Backup	_SATURN_KB/ compressed + encrypted	Weekly or after major updates	Google Drive or Backblaze B2 (\$6/mo)
Version Control	Key proposals, templates, SOPs	After every significant edit	Notion version history
System Snapshot	Full PC image before major installs	Before big upgrades	Macrium Reflect Free

What NOT to Automate

- Sending client emails — Saturn drafts, you review and send. Always.
- Signing or accepting contracts — never automate legal document approval.
- Financial transactions — never automate payments or invoice approvals in V1–2.
- Deleting files — build automations that archive, never permanently delete.
- Social media posting — AI content needs review before publishing for a professional brand.

11 Step-by-Step Build Plan

Phased Roadmap

Phase	Goal	Tools to Install	Cost	Time	Success Milestone
Phase 0 Planning	Organize files. Create _SATURN_KB/ folder. Define naming convention.	Everything (voidtools), DocFetcher	\$0	1 weekend 8–15 hrs	Clean _SATURN_KB/ with 10+ key documents organized and renamed.
Phase 1 Local Chat + KB	Install Ollama + AnythingLLM. Download Mistral 7B. Feed knowledge base. Test document Q&A.	Ollama, AnythingLLM, Mistral 7B, nomic-embed-text	\$0–15/mo	3–5 evenings 6–10 hrs	Ask Saturn: "What is my standard rate?" and get a correct answer from your files.
Phase 2 Voice + Automation	Add voice input (Whisper), TTS (Kokoro), Docker, and n8n. Build folder creator, weekly review, follow-up reminder.	Docker, WSL2, n8n, faster-whisper, Kokoro TTS, USB mic	\$50–90 hardware	4–8 evenings 10–20 hrs	Speak to Saturn and hear it respond. n8n creates folders and sends weekly summary.
Phase 3 Business Ops	CRM tracking, invoice follow-up, proposal versioning, auto-generated project briefs, Notion + Google Workspace integration.	n8n workflows, Notion API, Google Workspace API, Open WebUI	\$0–30/mo	6–12 evenings 15–30 hrs	Saturn handles weekly review automatically and drafts follow-up emails.
Phase 4 Server + Scale	Move Saturn to a dedicated mini-server for always-on access. Upgrade GPU for larger models. Optional fine-tuning.	Dedicated server, GPU upgrade, Ollama on server	\$500–1,500 hardware	2–4 weekends 10–20 hrs	Saturn accessible from phone, tablet, and workstation simultaneously.

Phase 1 Detailed Setup Order

1. Enable WSL2 in Windows: open PowerShell as admin, run: wsl --install
2. Download and install Docker Desktop from docker.com/products/docker-desktop. Restart PC.
3. Install Ollama from ollama.com. Open Command Prompt.
4. Download Mistral 7B: run `ollama pull mistral` (downloads ~4 GB, 5–15 min).
5. Download embedding model: run `ollama pull nomic-embed-text`
6. Test Ollama: run `ollama run mistral` and type "Hello" — confirm it responds.

7. Download and install AnythingLLM from anythingllm.com.
8. Open AnythingLLM  Settings  LLM Provider  select Ollama  choose mistral model.
9. Settings  Embedding Provider  select Ollama  choose nomic-embed-text.
10. Create a Workspace called "New Era Designs."
11. Drag _SATURN_KB/ folder contents into the workspace. Wait for indexing to complete (10–30 min).
12. Test: Ask "What is my standard hourly rate?" — verify Saturn returns the answer from your pricing doc.
13. Optional: add OpenAI API key in AnythingLLM Settings for cloud writing fallback.

12

Learning Path

What You Need to Learn — In Order

You are a technically capable creative professional, not a programmer. Learn the minimum needed for each phase, then move forward.

Concept	Plain English Explanation	Learn When	Difficulty
What a local LLM is	An AI model that runs on your own computer, not the cloud. Like a mini-ChatGPT on your hard drive.	Before Phase 1	Easy
What Ollama does	A program that downloads and runs AI models, similar to how Steam downloads games.	Phase 1 — Day 1	Very Easy
What RAG means	Saturn searches your documents first, then uses AI to answer. It's not memorizing your files — it's searching them on demand.	Phase 1	Easy
What a vector database does	Converts text into numbers that can be searched by meaning, not just keywords. "Rate card" and "pricing sheet" find the same result.	Phase 1 (concept only — AnythingLLM handles it)	Easy — no action needed
What Docker is	Runs apps in isolated "containers" inside your computer. Required for Open WebUI, n8n server, Kokoro TTS in Phase 2.	Before Phase 2	Medium
Basic Docker commands	docker pull (download), docker run (start), docker ps (list), docker stop (stop). Five commands cover 90% of what you need.	Phase 2	Easy once understood
What an API is	A way for two programs to talk to each other. An API key is a password that lets AnythingLLM access OpenAI.	Phase 1 (optional)	Very Easy
What n8n does	A visual automation tool where you connect blocks on a canvas to build workflows. No coding for basic tasks.	Phase 2	Medium
What an embedding model is	A smaller AI model (nomic-embed-text) that converts text to vectors. Required for document search.	Phase 1 — Day 1	Easy — 1 install command
What a system prompt is	Instructions at the start of every session: "You are Saturn, the AI assistant for New Era Designs." Defines Saturn's role and behavior.	Phase 1 — after first test	Very Easy
What fine-tuning is	Retraining a model on your own data. Complex, time-consuming. You do NOT need this for V1-3.	Phase 4 — optional	High — ignore for now

What to Ignore Completely Right Now

- Training your own AI model from scratch — not practical on consumer hardware.
- LangChain and LlamaIndex Python frameworks — AnythingLLM handles this visually.
- Kubernetes and cloud deployment — enterprise only, not relevant for a solo business.
- Machine learning theory — interesting but not actionable for your use case.

- Building a custom chat interface from code — Open WebUI or AnythingLLM handle this.

13

Final Recommendation

Simplest Realistic Saturn V1 Stack — Start This Weekend

#	Tool	Role in Saturn	Where to Get It
1	Ollama	Model runner — the engine behind everything	ollama.com
2	Mistral 7B via Ollama	Primary local AI model for all tasks	Run: ollama pull mistral
3	nomic-embed-text via Ollama	Document embedding for file search	Run: ollama pull nomic-embed-text
4	AnythingLLM (Desktop)	Saturn interface + document Q&A; system	anythingllm.com
5	OpenAI API key (optional)	Cloud fallback for complex writing	platform.openai.com
6	Everything (voidtools)	Instant file name search across all drives	voidtools.com
7	DocFetcher	Search text inside PDFs and DOCX files	docfetcher.sourceforge.net
8	Obsidian	Local knowledge base (SOPs, naming rules)	obsidian.md
9	_SATURN_KB/ folder	Your organized knowledge base — root of Saturn	Create manually on your drive

Total setup time: 3–5 hours across one evening and one morning. All software is free. Optional OpenAI API: ~\$10/mo for typical solo use. Your hardware is already sufficient.

Best Long-Term Architecture (6–12 Months)

Component	Long-Term Tool	Why
Local model runner	Ollama	Stable, GPU-accelerated, widely supported.
Main interface	Open WebUI (Docker)	More features than AnythingLLM once comfortable with Docker.
Automation engine	n8n self-hosted (Docker)	Most powerful self-hosted automation. 400+ native integrations.
Voice input	faster-whisper + hotkey script	Local, private, accurate, GPU-accelerated.
Voice output	Kokoro TTS (Docker)	Best local TTS quality in 2026. 54 voice options.
Local model (upgraded)	Mistral 22B or Qwen2.5 32B (needs 16–24 GB VRAM)	Near-GPT-4 quality locally if you upgrade GPU.
Cloud writing fallback	OpenAI GPT-4.1 or Claude Sonnet (API, drafting only)	Premium quality for proposals and contracts.
CRM + project tracking	Notion (existing) + n8n integrations	Notion is already in your stack. n8n connects everything.

Component	Long-Term Tool	Why
Always-on server (Phase 4)	Beelink SER7 mini PC (~\$350)	Runs Saturn 24/7 without tying up your workstation.

REF

Recommended Saturn V1 Stack

Ollama	ollama.com	OpenAI API	platform.openai.com
Mistral 7B	ollama pull mistral	Everything	voidtools.com
nomic-embed-text	ollama pull nomic-embed-text	DocFetcher	docfetcher.sourceforge.net
AnythingLLM	anythingllm.com	Obsidian	obsidian.md

Saturn System Prompt — Paste Into AnythingLLM

You are Saturn, the private AI assistant for New Era Designs, LLC. Your role is to help Chris — a solo brand identity, wayfinding, and environmental graphics designer — run his business more efficiently.

Always answer based on the knowledge base documents first. If the answer is not in the documents, say so clearly and offer your best general guidance.

Be direct, professional, and practical. Use plain English. Avoid excessive disclaimers. Format responses clearly with bullet points or numbered steps when appropriate.

You help with: proposal drafting, pricing and scope estimates, project summaries, naming conventions, meeting notes, file search, creative brief extraction, client follow-up drafts, and business operations. You do not handle legal advice, financial transactions, or client communications without human review.

7DAY First 7 Days Setup Checklist

Follow this checklist in order. Do not skip ahead. Total time: 8–12 hours spread across one week.

Day 1: File Cleanup (2–3 hours)

1. Download and install Everything (voidtools.com). Search for duplicate filenames and resolve them.
2. Create the _SATURN_KB/ folder structure on your external drive (see Section 7).
3. Move top 10–15 key documents (proposals, pricing, SOPs) into _SATURN_KB/ with correct naming.
4. Install DocFetcher. Index _SATURN_KB/ and confirm you can search text inside your files.

Day 2: Install Core Saturn Tools (1–2 hours)

5. Download and install Ollama from ollama.com. Open Command Prompt. Run: ollama --version
6. Download Mistral 7B: run ollama pull mistral (downloads ~4 GB, takes 5–15 min).
7. Download embedding model: run ollama pull nomic-embed-text
8. Test Ollama: run ollama run mistral — type "Hello" and confirm it responds.
9. Download and install AnythingLLM Desktop from anythingllm.com.

Day 3: Configure Saturn + Load Knowledge Base (2–3 hours)

10. Open AnythingLLM ≡ Settings ≡ LLM Provider ≡ select Ollama ≡ choose mistral.
11. Settings ≡ Embedding Provider ≡ select Ollama ≡ choose nomic-embed-text.
12. Create a Workspace called "New Era Designs."
13. Drag _SATURN_KB/ files into the workspace. Wait for indexing (10–30 min for first batch).
14. Optional: go to Settings ≡ add your OpenAI API key for cloud writing fallback.

Day 4: System Prompt + First Tests (1 hour)

15. Open AnythingLLM workspace settings ≡ paste the System Prompt from the reference page.
16. Test 1: Ask "What is my standard hourly design rate?" — verify Saturn returns your rate.
17. Test 2: Ask "Summarize the most recent proposal." — verify it reads your proposal correctly.
18. Test 3: Paste a sample client email and ask Saturn to draft a professional response.
19. Test 4: Paste raw meeting notes and ask for a structured summary with action items.

Day 5: Install Obsidian + Build Notes (1 hour)

20. Download and install Obsidian from obsidian.md.
21. Create a vault pointing to your _SATURN_KB/ folder so notes are stored as plain .md files.
22. Create your first note: NERA_NAMING_RULES.md — document your file naming convention.
23. Create a second note: NERA_PRICING_LOGIC.md — write out your rates and discount rules.
24. Re-import these notes into AnythingLLM so Saturn knows your naming and pricing logic.

Day 6: Real Work Test (1 hour)

25. Use Saturn for one real task today: draft a follow-up email, generate a scope estimate, or summarize a project.
26. Note what Saturn got right and what it missed. Add missing information to _SATURN_KB/ to improve it.
27. Identify the 3 most useful things Saturn did today.

Day 7: Plan Phase 2 + Mark Complete (30 min)

28. Review Phase 2 shopping list: USB microphone (\$50–70) and Docker Desktop (free).
29. Schedule a "Saturn Phase 2 Build Day" on your calendar 3–4 weeks from now.
30. Write a short note: what Saturn does well now, what you want it to do in Phase 2.
31. Saturn V1 is live. The business intelligence grows every time you add new documents.

Saturn V1 is now live. From here, it only gets more useful as you feed it more documents and build out more workflows. The foundation is solid. The business intelligence grows over time.